

AAYUSH SAXENA

+91 - 9672121399

Gurugram, Haryana

m.aayushsaxena@gmail.com

Forward Deployed AI Engineer

github.com/storytellingengineer/
linkedin.com/in/storytellingengineer/

AI Engineer with 6+ years of experience building, deploying, and validating AI systems across enterprise environments. Experienced in LLM applications, RAG, AI agents, evaluation frameworks, and customer-facing AI solutions.

EDUCATION

- **M.Tech** in Applied Artificial Intelligence
[Visvesvaraya National Institute of Technology, Nagpur](#)
Dec 2023 – Dec 2025
 - **Dissertation:** *Multimodal Stress Detection using Deep Learning & Sensor Fusion (StressXNet)*
- **B.Tech** in Computer Science Engineering
[JECRC University, Jaipur](#)
May 2016 - 2020

FEATURED PROJECTS

AI Quality Analysis & Validation Framework

Tech stack: Langfuse, LLM-as-a-Judge, Responsible AI

- Designed an enterprise AI Validation framework for production applications covering relevance, prompt injection, hallucination & PII Detection.
- Built custom evaluator prompts and automated workflows to improve AI reliability.
- Performed root cause analysis for LLM failures, retrieval issues and guardrail violations.

Retail Intelligence Assistant (RAG + AI Agents)

Tech stack: LangGraph, FastAPI, Milvus, Vertex AI

- Built a production RAG assistant integrating enterprise knowledge, vector search and tool-calling workflows.
- Developed API-based workflows and optimized retrieval and response generation.

TECHNICAL SKILLS

Languages – Python, R, SQL

AI System – LLMs, RAG, Agentic AI, AI Evaluation
Responsible AI, Prompt Engineering, GenAI

ML - Machine Learning, Deep Learning, Statistics

Engineering - CI/CD, Docker, Git, FastAPI, REST APIs

Cloud & Data - VertexAI, Milvus, Qdrant, FAISS

CORE COMPETENCIES

Generative AI — LLM, RAG, AI Agents, Tool Calling

AI Reliability — Evaluation, LLM-as-a-Judge, Guardrails,
Hallucination Detection

AI Engineering — APIs, Retrieval Pipelines, Vector
Databases, Deployment

Enterprise AI — Solution Design, Validation, Governance,
Stakeholder Collaboration

WORK EXPERIENCE

Assistant Manager - AI Validation & Testing

[EY GDS](#)

March 2026 – Present

Gurugram, India

- Designed validation strategies for enterprise-scale LLM applications across GxP/Non-GxP regulated pharmaceutical and manufacturing workflows.
- Built automated AI evaluation pipelines using Langfuse, LLM-as-a-Judge, and custom evaluator covering response quality, groundedness, hallucination, PII, and toxicity detection.
- Partnered with client, product and engineering teams to troubleshoot production issues through root cause analysis and validation testing.
- Improved enterprise AI reliability by identifying guardrail failures, retrieval issues, prompt attacks and evaluation gaps before production release.
- Defined AI governance documentation, validation evidence, and release timeline for timely deployment.

Associate Data Scientist – GenAI

July 2025 – Nov 2025

[Easyrewardz Software Services](#)

Gurugram, India

- Built production RAG application integrating vector search, tool-calling, and enterprise knowledge bases using LangChain/LangGraph.
- Improved response accuracy by **18%** and reduced latency by **~30%** using retrieval tuning.

Data Scientist – Predictive ML

Feb 2023 – July 2025

[Aureus Analytics, a Gradatim Insurance Tech Company](#)

Mumbai, India

- Built 10+ predictive ML models for insurance fraud detection, claim prediction and cross-sell.
- Improved performance by **~20%** using feature engineering and ensemble methods.
- Translated ML outputs into actionable insights by partnering with business and cross-functional teams.

Associate Data Analyst

Aug 2021 – May 2022

[GlobalLogic \(Hitachi Group Company\)](#)

Gurugram, India

- Led A/B testing and statistical analysis
- Built dashboards and reporting pipelines to support decision-making.

CO-CURRICULUM ACTIVITIES

- **Internships during College (2016–2020)** – Gained hands-on experience in diverse roles including ML, Web development, E-summit organization, and lead generation.
- **Tech Blogger & Mentor** – Published practical content on applied ML and mentored other data scientists via GitHub/LinkedIn.